

rock-vs-mine-prediction

0.0.1 Importing the Dependencies

```
[70]: # import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

0.0.2 Data Collection and Data Processing

```
[71]: # loading the dataframe
df=pd.read_csv('sonar data.csv',header=None)
```

```
[11]: df.head()
```

```
[11]:
```

	0	1	2	3	4	5	6	7	8	\
0	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601	0.3109	
1	0.0453	0.0523	0.0843	0.0689	0.1183	0.2583	0.2156	0.3481	0.3337	
2	0.0262	0.0582	0.1099	0.1083	0.0974	0.2280	0.2431	0.3771	0.5598	
3	0.0100	0.0171	0.0623	0.0205	0.0205	0.0368	0.1098	0.1276	0.0598	
4	0.0762	0.0666	0.0481	0.0394	0.0590	0.0649	0.1209	0.2467	0.3564	

	9	...	51	52	53	54	55	56	57	\
0	0.2111	...	0.0027	0.0065	0.0159	0.0072	0.0167	0.0180	0.0084	
1	0.2872	...	0.0084	0.0089	0.0048	0.0094	0.0191	0.0140	0.0049	
2	0.6194	...	0.0232	0.0166	0.0095	0.0180	0.0244	0.0316	0.0164	
3	0.1264	...	0.0121	0.0036	0.0150	0.0085	0.0073	0.0050	0.0044	
4	0.4459	...	0.0031	0.0054	0.0105	0.0110	0.0015	0.0072	0.0048	

	58	59	60
0	0.0090	0.0032	R
1	0.0052	0.0044	R
2	0.0095	0.0078	R
3	0.0040	0.0117	R

```
4  0.0107  0.0094   R
```

```
[5 rows x 61 columns]
```

```
[12]: # number of rows and columns
df.shape
```

```
[12]: (208, 61)
```

```
[13]: # info
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 208 entries, 0 to 207
```

```
Data columns (total 61 columns):
```

#	Column	Non-Null Count	Dtype
0	0	208 non-null	float64
1	1	208 non-null	float64
2	2	208 non-null	float64
3	3	208 non-null	float64
4	4	208 non-null	float64
5	5	208 non-null	float64
6	6	208 non-null	float64
7	7	208 non-null	float64
8	8	208 non-null	float64
9	9	208 non-null	float64
10	10	208 non-null	float64
11	11	208 non-null	float64
12	12	208 non-null	float64
13	13	208 non-null	float64
14	14	208 non-null	float64
15	15	208 non-null	float64
16	16	208 non-null	float64
17	17	208 non-null	float64
18	18	208 non-null	float64
19	19	208 non-null	float64
20	20	208 non-null	float64
21	21	208 non-null	float64
22	22	208 non-null	float64
23	23	208 non-null	float64
24	24	208 non-null	float64
25	25	208 non-null	float64
26	26	208 non-null	float64
27	27	208 non-null	float64
28	28	208 non-null	float64
29	29	208 non-null	float64

```

30 30      208 non-null    float64
31 31      208 non-null    float64
32 32      208 non-null    float64
33 33      208 non-null    float64
34 34      208 non-null    float64
35 35      208 non-null    float64
36 36      208 non-null    float64
37 37      208 non-null    float64
38 38      208 non-null    float64
39 39      208 non-null    float64
40 40      208 non-null    float64
41 41      208 non-null    float64
42 42      208 non-null    float64
43 43      208 non-null    float64
44 44      208 non-null    float64
45 45      208 non-null    float64
46 46      208 non-null    float64
47 47      208 non-null    float64
48 48      208 non-null    float64
49 49      208 non-null    float64
50 50      208 non-null    float64
51 51      208 non-null    float64
52 52      208 non-null    float64
53 53      208 non-null    float64
54 54      208 non-null    float64
55 55      208 non-null    float64
56 56      208 non-null    float64
57 57      208 non-null    float64
58 58      208 non-null    float64
59 59      208 non-null    float64
60 60      208 non-null    object
dtypes: float64(60), object(1)
memory usage: 99.3+ KB

```

```
[14]: # describe // statistical measures of the data
df.describe()
```

```
[14]:
```

	0	1	2	3	4	5	\
count	208.000000	208.000000	208.000000	208.000000	208.000000	208.000000	
mean	0.029164	0.038437	0.043832	0.053892	0.075202	0.104570	
std	0.022991	0.032960	0.038428	0.046528	0.055552	0.059105	
min	0.001500	0.000600	0.001500	0.005800	0.006700	0.010200	
25%	0.013350	0.016450	0.018950	0.024375	0.038050	0.067025	
50%	0.022800	0.030800	0.034300	0.044050	0.062500	0.092150	
75%	0.035550	0.047950	0.057950	0.064500	0.100275	0.134125	
max	0.137100	0.233900	0.305900	0.426400	0.401000	0.382300	

	6	7	8	9	...	50	\
count	208.000000	208.000000	208.000000	208.000000	...	208.000000	
mean	0.121747	0.134799	0.178003	0.208259	...	0.016069	
std	0.061788	0.085152	0.118387	0.134416	...	0.012008	
min	0.003300	0.005500	0.007500	0.011300	...	0.000000	
25%	0.080900	0.080425	0.097025	0.111275	...	0.008425	
50%	0.106950	0.112100	0.152250	0.182400	...	0.013900	
75%	0.154000	0.169600	0.233425	0.268700	...	0.020825	
max	0.372900	0.459000	0.682800	0.710600	...	0.100400	

	51	52	53	54	55	56	\
count	208.000000	208.000000	208.000000	208.000000	208.000000	208.000000	
mean	0.013420	0.010709	0.010941	0.009290	0.008222	0.007820	
std	0.009634	0.007060	0.007301	0.007088	0.005736	0.005785	
min	0.000800	0.000500	0.001000	0.000600	0.000400	0.000300	
25%	0.007275	0.005075	0.005375	0.004150	0.004400	0.003700	
50%	0.011400	0.009550	0.009300	0.007500	0.006850	0.005950	
75%	0.016725	0.014900	0.014500	0.012100	0.010575	0.010425	
max	0.070900	0.039000	0.035200	0.044700	0.039400	0.035500	

	57	58	59
count	208.000000	208.000000	208.000000
mean	0.007949	0.007941	0.006507
std	0.006470	0.006181	0.005031
min	0.000300	0.000100	0.000600
25%	0.003600	0.003675	0.003100
50%	0.005800	0.006400	0.005300
75%	0.010350	0.010325	0.008525
max	0.044000	0.036400	0.043900


[8 rows x 60 columns]

```
[17]: # Count of Mine (M) and Rock (R)
df[60].value_counts()
```

```
[17]: 60
M      111
R       97
Name: count, dtype: int64
```

```
[79]: # count plot
plt.figure(figsize=(3,2))
sns.countplot(x=df[60], data=df)
```

```
[79]: <Axes: xlabel='60', ylabel='count'>
```



```
[80]: # check null vlaues
df.isnull().sum()
```

```
[80]: 0      0
      1      0
      2      0
      3      0
      4      0
      ..
      56     0
      57     0
      58     0
      59     0
      60     0
      Length: 61, dtype: int64
```

```
[25]: # get mean of both Mine and Rock which is M and R
df.groupby(60).mean()
```

```
[25]:
```

	0	1	2	3	4	5	6	\
60								
M	0.034989	0.045544	0.050720	0.064768	0.086715	0.111864	0.128359	
R	0.022498	0.030303	0.035951	0.041447	0.062028	0.096224	0.114180	
	7	8	9	...	50	51	52	53 \
60				...				
M	0.149832	0.213492	0.251022	...	0.019352	0.016014	0.011643	0.012185
R	0.117596	0.137392	0.159325	...	0.012311	0.010453	0.009640	0.009518
	54	55	56	57	58	59		
60								
M	0.009923	0.008914	0.007825	0.009060	0.008695	0.006930		

```
R    0.008567  0.007430  0.007814  0.006677  0.007078  0.006024
```

```
[2 rows x 60 columns]
```

0.0.3 Separating data and labels

```
[27]: X= df.drop(columns=60, axis=1)
      Y= df[60]
```

```
[28]: print (X)
      print (Y)
```

	0	1	2	3	4	5	6	7	8	\
0	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601	0.3109	
1	0.0453	0.0523	0.0843	0.0689	0.1183	0.2583	0.2156	0.3481	0.3337	
2	0.0262	0.0582	0.1099	0.1083	0.0974	0.2280	0.2431	0.3771	0.5598	
3	0.0100	0.0171	0.0623	0.0205	0.0205	0.0368	0.1098	0.1276	0.0598	
4	0.0762	0.0666	0.0481	0.0394	0.0590	0.0649	0.1209	0.2467	0.3564	
..	
203	0.0187	0.0346	0.0168	0.0177	0.0393	0.1630	0.2028	0.1694	0.2328	
204	0.0323	0.0101	0.0298	0.0564	0.0760	0.0958	0.0990	0.1018	0.1030	
205	0.0522	0.0437	0.0180	0.0292	0.0351	0.1171	0.1257	0.1178	0.1258	
206	0.0303	0.0353	0.0490	0.0608	0.0167	0.1354	0.1465	0.1123	0.1945	
207	0.0260	0.0363	0.0136	0.0272	0.0214	0.0338	0.0655	0.1400	0.1843	

	9	...	50	51	52	53	54	55	56	\
0	0.2111	...	0.0232	0.0027	0.0065	0.0159	0.0072	0.0167	0.0180	
1	0.2872	...	0.0125	0.0084	0.0089	0.0048	0.0094	0.0191	0.0140	
2	0.6194	...	0.0033	0.0232	0.0166	0.0095	0.0180	0.0244	0.0316	
3	0.1264	...	0.0241	0.0121	0.0036	0.0150	0.0085	0.0073	0.0050	
4	0.4459	...	0.0156	0.0031	0.0054	0.0105	0.0110	0.0015	0.0072	
..	
203	0.2684	...	0.0203	0.0116	0.0098	0.0199	0.0033	0.0101	0.0065	
204	0.2154	...	0.0051	0.0061	0.0093	0.0135	0.0063	0.0063	0.0034	
205	0.2529	...	0.0155	0.0160	0.0029	0.0051	0.0062	0.0089	0.0140	
206	0.2354	...	0.0042	0.0086	0.0046	0.0126	0.0036	0.0035	0.0034	
207	0.2354	...	0.0181	0.0146	0.0129	0.0047	0.0039	0.0061	0.0040	

	57	58	59
0	0.0084	0.0090	0.0032
1	0.0049	0.0052	0.0044
2	0.0164	0.0095	0.0078
3	0.0044	0.0040	0.0117
4	0.0048	0.0107	0.0094
..
203	0.0115	0.0193	0.0157
204	0.0032	0.0062	0.0067

```

205  0.0138  0.0077  0.0031
206  0.0079  0.0036  0.0048
207  0.0036  0.0061  0.0115

```

```

[208 rows x 60 columns]
0      R
1      R
2      R
3      R
4      R
..
203    M
204    M
205    M
206    M
207    M
Name: 60, Length: 208, dtype: object

```

0.0.4 Train and Test data

```
[37]: X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size = 0.1, stratify_
      ↪= Y, random_state = 1)
```

```
[39]: print(X.shape, X_train.shape, X_test.shape)
```

```
(208, 60) (187, 60) (21, 60)
```

```
[39]: print(X.shape, X_train.shape, X_test.shape)
```

```
(208, 60) (187, 60) (21, 60)
```

1 Model Training >> Logistic Regression

```
[44]: model = LogisticRegression()
```

```
[45]: # training the logistic Regrsson model with training data
      model.fit(X_train, Y_train)
```

```
[45]: LogisticRegression()
```

1.0.1 Model Evaluation

```
[68]: # prediction
      X_train_prediction = model.predict(X_train)
```

```
[50]: # accuarcy of training data
      training_data_accuarcy= accuracy_score(X_train_prediction, Y_train)
```

```
print(training_data_accuarcy)
```

0.8342245989304813

```
[66]: # compare the prediction value with Actual value
result_df=pd.DataFrame({'Actual': Y_train, 'Prediction':X_train_prediction})
```

```
[69]: # get output and check accuracy
print(result_df.tail(10))
print(f"\n Accuracy on the test set: {training_data_accuarcy}")
```

	Actual	Prediction
177	M	M
171	M	M
94	R	R
185	M	M
67	R	R
140	M	M
5	R	R
154	M	M
131	M	M
203	M	M

Accuracy on the test set: 0.8342245989304813

2 Making Predicting data

```
[65]: input_data = (0.0526,0.0563,0.1219,0.1206,0.0246,0.1022,0.0539,0.0439,0.2291,0.
↪1632,0.2544,0.2807,0.3011,0.3361,0.3024,0.2285,0.2910,0.1316,0.1151,0.3404,0.
↪5562,0.6379,0.6553,0.7384,0.6534,0.5423,0.6877,0.7325,0.7726,0.8229,0.8787,0.
↪9108,0.6705,0.6092,0.7505,0.4775,0.1666,0.3749,0.3776,0.2106,0.5886,0.5628,0.
↪2577,0.5245,0.6149,0.5123,0.3385,0.1499,0.0546,0.0270,0.0380,0.0339,0.0149,0.
↪0335,0.0376,0.0174,0.0132,0.0103,0.0364,0.0208)

# Changing the input_data to a numpy array
input_data_as_numpy_array = np.asarray(input_data)

# Reshape the np array as we are predicting for one instance
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)
prediction = model.predict(input_data_reshaped)
print(prediction)

if(prediction[0]=='R'):
    print("The Object is a Rock")
else:
    print("The Object is a Mine")
```


['M']

The Object is a Mine